Operating Systems

Memory Management

Operating Systems

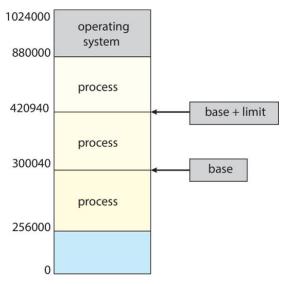
Main Memory

Background

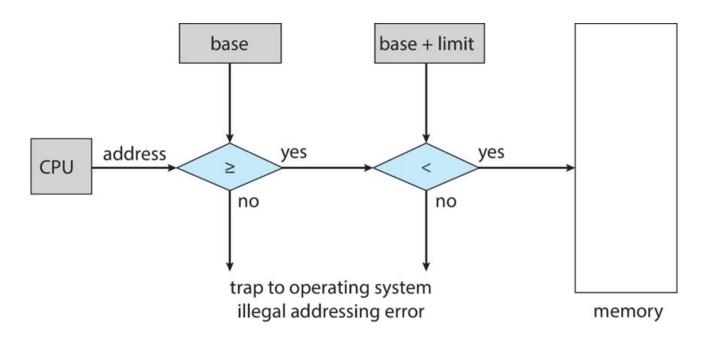
- Program must be brought (from disk) into memory and placed within a process for it to be run
- ☐ Main memory and registers are only storage CPU can access directly
- ☐ Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- Register access in one CPU clock (or less)
- Main memory can take many cycles, causing a **stall**, since it does not have the data required to complete the instruction that it is executing
- Cache sits between main memory and CPU registers for fast access
- ☐ Protection of memory required to ensure correct operation

Base and Limit Registers

- A pair of base and limit registers define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



Hardware Address Protection



Operating Systems **Address Space**

Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate physical address space is central to proper memory management
 - Logical address generated by the CPU; also referred to as virtual address
 - Physical address address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- Logical address space is the set of all logical addresses generated by a program
- Physical address space is the set of all physical addresses generated by a program

Operating Systems

Paging

Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called frames
 - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size N pages, need to find N free frames and load program
- Set up a page table to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation

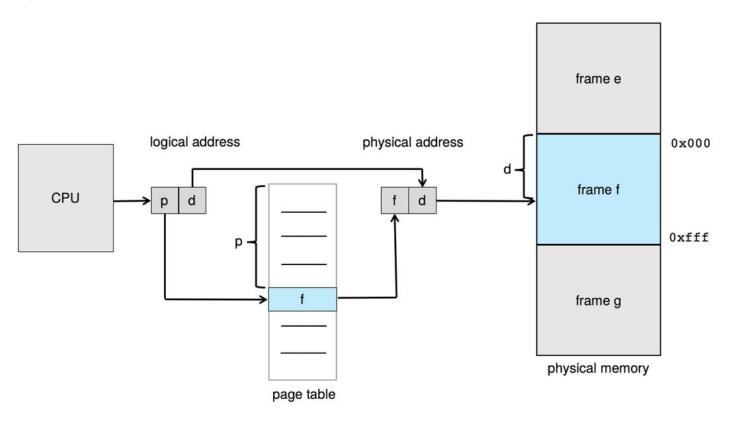
Address Translation Scheme

- Address generated by CPU is divided into:
 - Page number (p) used as an index into a page table which contains base address of each page in physical memory
 - Page offset (*d*) combined with base address to define the physical memory address that is sent to the memory unit

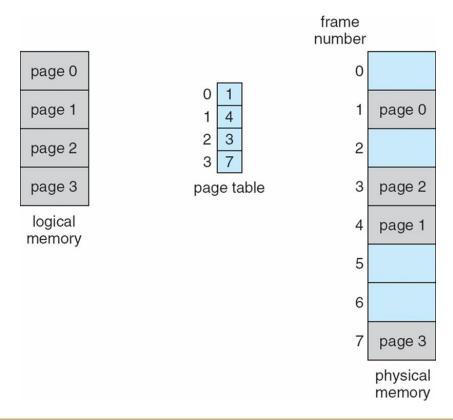
page number	page offset
р	d
m -n	n

• For given logical address space 2^m and page size 2^n

Paging Hardware



Paging Model of Logical and Physical Memory



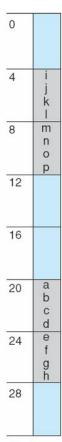
Paging Example

n=2 and *m*=432-byte memory and 4-byte pages

0	a
1	b
2	С
3	d
4 5	е
5	f
6	g
7	g h
8	i
9	j k
10	k
11	1
12	m
13	n
14	0
15	р

1	6	
2	1	
3	2	
oage	ta	ble

0 5



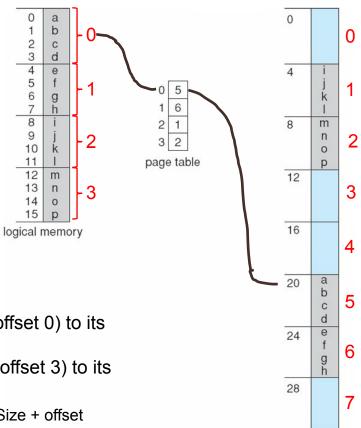
physical memory

Paging Example

- m is used to determine the number of logical addresses, 2^m = 16
- n indicates the offset within each logical address

$$n=2$$
 and $m=4$

32-byte memory and 4-byte pages



- Map logical address 0 (page 0, offset 0) to its corresponding physical address
- Map Logical address 3 (page 0, offset 3) to its corresponding physical address

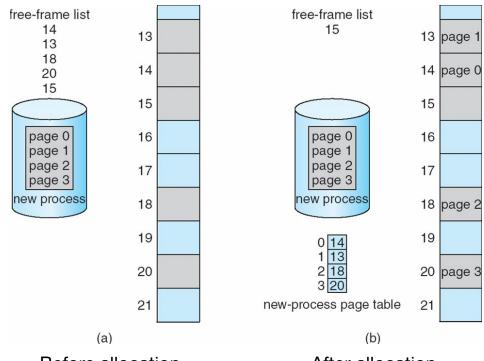
frameNumber * frameSize + offset

physical memory

Size of Page

- Calculating internal fragmentation
 - \circ Page size = 2,048 bytes
 - \circ Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - \circ Internal fragmentation of 2,048 1,086 = 962 bytes
 - \circ Worst case fragmentation = 1 frame 1 byte
 - On average fragmentation = 1 / 2 frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes growing over time
 - Solaris supports two page sizes 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory

Free Frames



Before allocation

After allocation

Operating Systems Implementation of Page Table

Implementation of Page Table

- Page table is kept in main memory
- Page-table base register (PTBR) points to the page table
- Page-table length register (PTLR) indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation lookaside buffers** (TLBs)

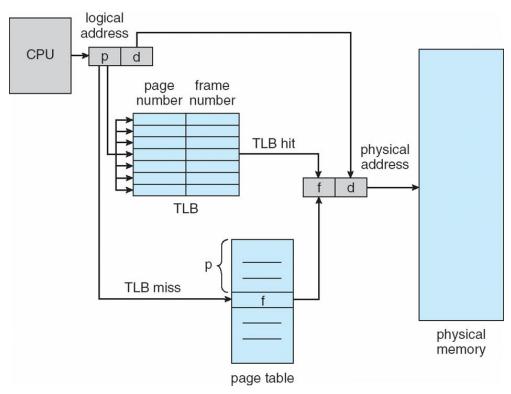
Associative Memory

• Associative memory – parallel search

Page #	Frame #

- Address translation (p, d)
 - \circ If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory

Paging Hardware With TLB



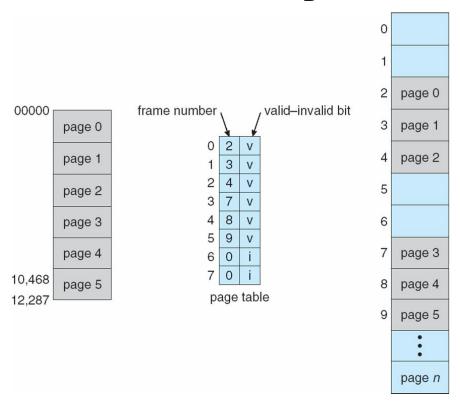
Effective Access Time

- Associative Lookup = ε time unit
 - Can be < 10% of memory access time
- Hit ratio = α
 - Hit ratio percentage of times that a page number is found in the associative registers;
- Effective Access Time (EAT):
- Consider $\alpha = 80\%$, $\epsilon = 20$ ns for TLB search, 100ns for memory access
 - \circ EAT = 0.80 x 120 + 0.20 x 220 = 140ns
- Consider more realistic hit ratio -> $\alpha = 99\%$, $\epsilon = 20$ ns for TLB search, 100ns for memory access
 - \circ EAT = 0.99 x 120 + 0.01 x 220 = 121ns

Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - "valid" indicates that the associated page is in the process' logical address space, and is thus a legal page
 - "invalid" indicates that the page is not in the process' logical address space
 - Or use page-table length register (PTLR)
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table



Operating Systems

Shared Pages

Shared Pages

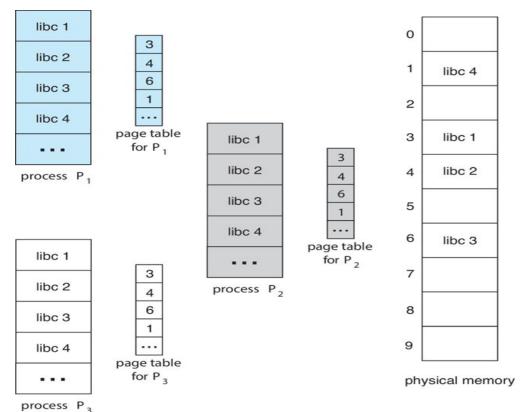
Shared code

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for inter-process communication if sharing of read-write pages is allowed

Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example



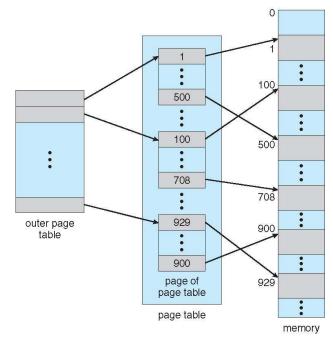
Operating Systems **Hierarchical Page Table**

Structure of the Page Table

- Memory structures for paging can get huge using straightforward methods
 - Consider a 32-bit logical address space as on modern computers
 - \circ Page size of 4 KB (2¹²)
 - Page table would have 1 million entries $(2^{32} / 2^{12})$
 - If each entry is 4 Bytes -> 4 MB of physical address space / memory for page table alone
 - That amount of memory used to cost a lot
 - Don't want to allocate that contiguously in main memory

Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table



Two-Level Paging Example

- A logical address (on 32-bit machine with 1KB page size) is divided into:
 - o a page number consisting of 22 bits
 - o a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - o a 12-bit page number
 - o a 10-bit page offset
- Thus, a logical address is as follows:

page number		page onset
p_1	p_2	d
12	10	10

- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table
- Known as forward-mapped page table

Operating Systems

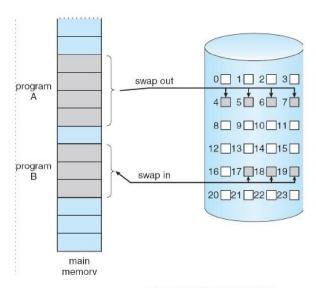
Virtual Memory

Background

- The term "virtual memory" refers to something which appears to be present but actually it is not.
- The virtual memory technique allows users to use more memory for a program than the real memory of a computer.
- Virtual memory is a **concept** that we use when we have processes that exceed the main memory.
- When computer runs out of physical memory, it writes its requirement to the hard disc in a swap file as "virtual memory".

Demand Paging

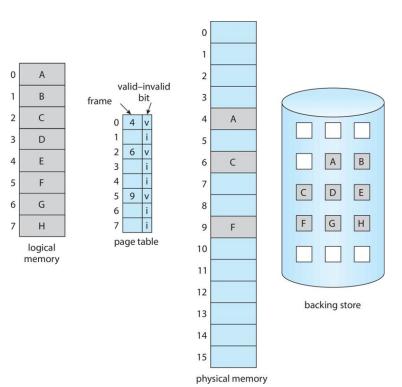
- ☐ Bring a page into memory only when it is needed
- Less I/O needed
- Less memory needed
- Faster response
- More users
- ☐ Disadvantage: Page fault interrupt
- ☐ Required hardware support:
- Page Table with valid-invalid bit
- Secondary memory



Demand Paging in OS

Valid-Invalid Bit

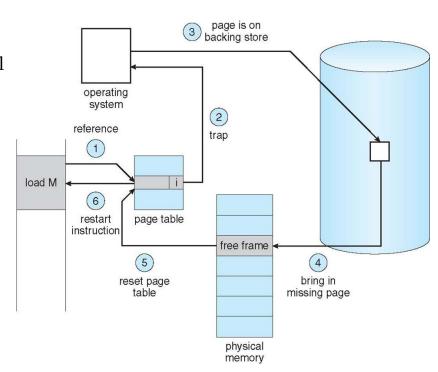
- An extra bit in the page table which indicates the existence of the page in the main memory.
- ☐ Attempt to access page
- ☐ If page is valid (in memory) then continue processing instruction as normal.
- ☐ If page is invalid then a page-fault trap / page-fault interrupt occurs.
- \square Page is needed \Rightarrow reference to it
- Invalid reference \Rightarrow abort
- Not-in-memory ⇒ bring to memory



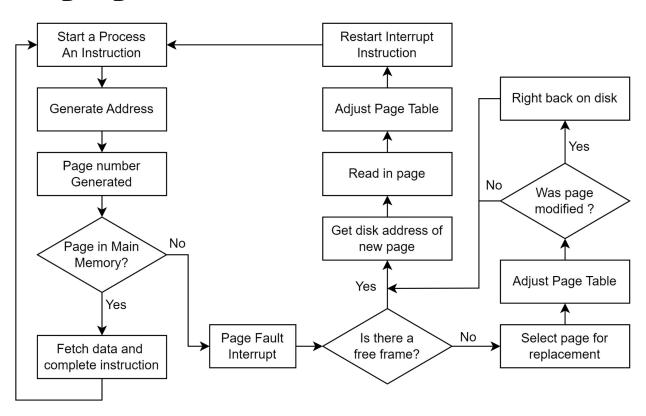
Page Fault

If there is ever a reference to a page, first reference will trap to $OS \Rightarrow page fault$

- 1. OS looks at another table to decide:
 - Invalid reference \Rightarrow abort.
 - Just not in memory.
- 2. Find empty/ free frame.
- 3. Load page from disk into frame.
- 4. Reset tables, validation bit = 1.
- 5. Restart instruction that caused page fault



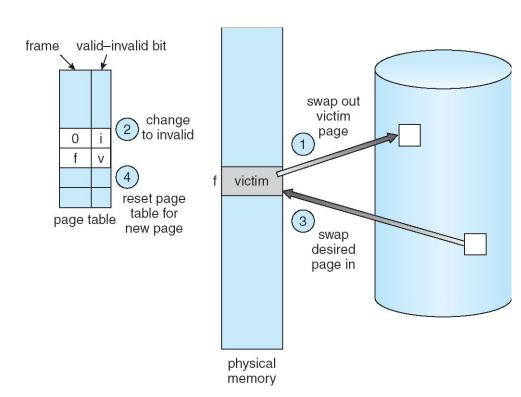
Demand Paging Flowchart



Page Replacement

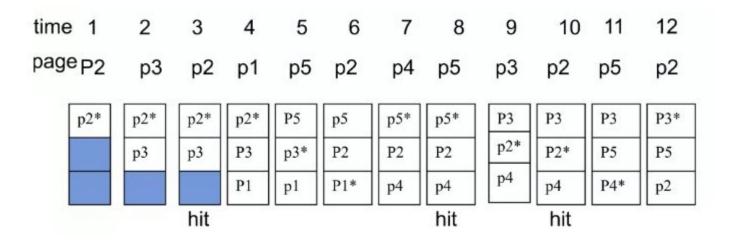
Page Replacement Algorithms:

- ☐ FIFO (First In First Out)
- ☐ LRU (Least Recently Used)
- □ OPT (Optimal)



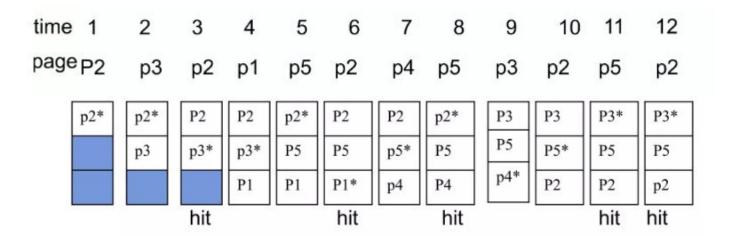
FIFO (First In First Out)

□ Selects the page for replacement that has been in the memory for the longest amount of time



LRU (Least Recently Used)

- Replace the least recently used page in the past
- Can be implemented by associating a counter with every page that is in main memory



Optimal

Replace the page which is not used in longest dimension of time in future

